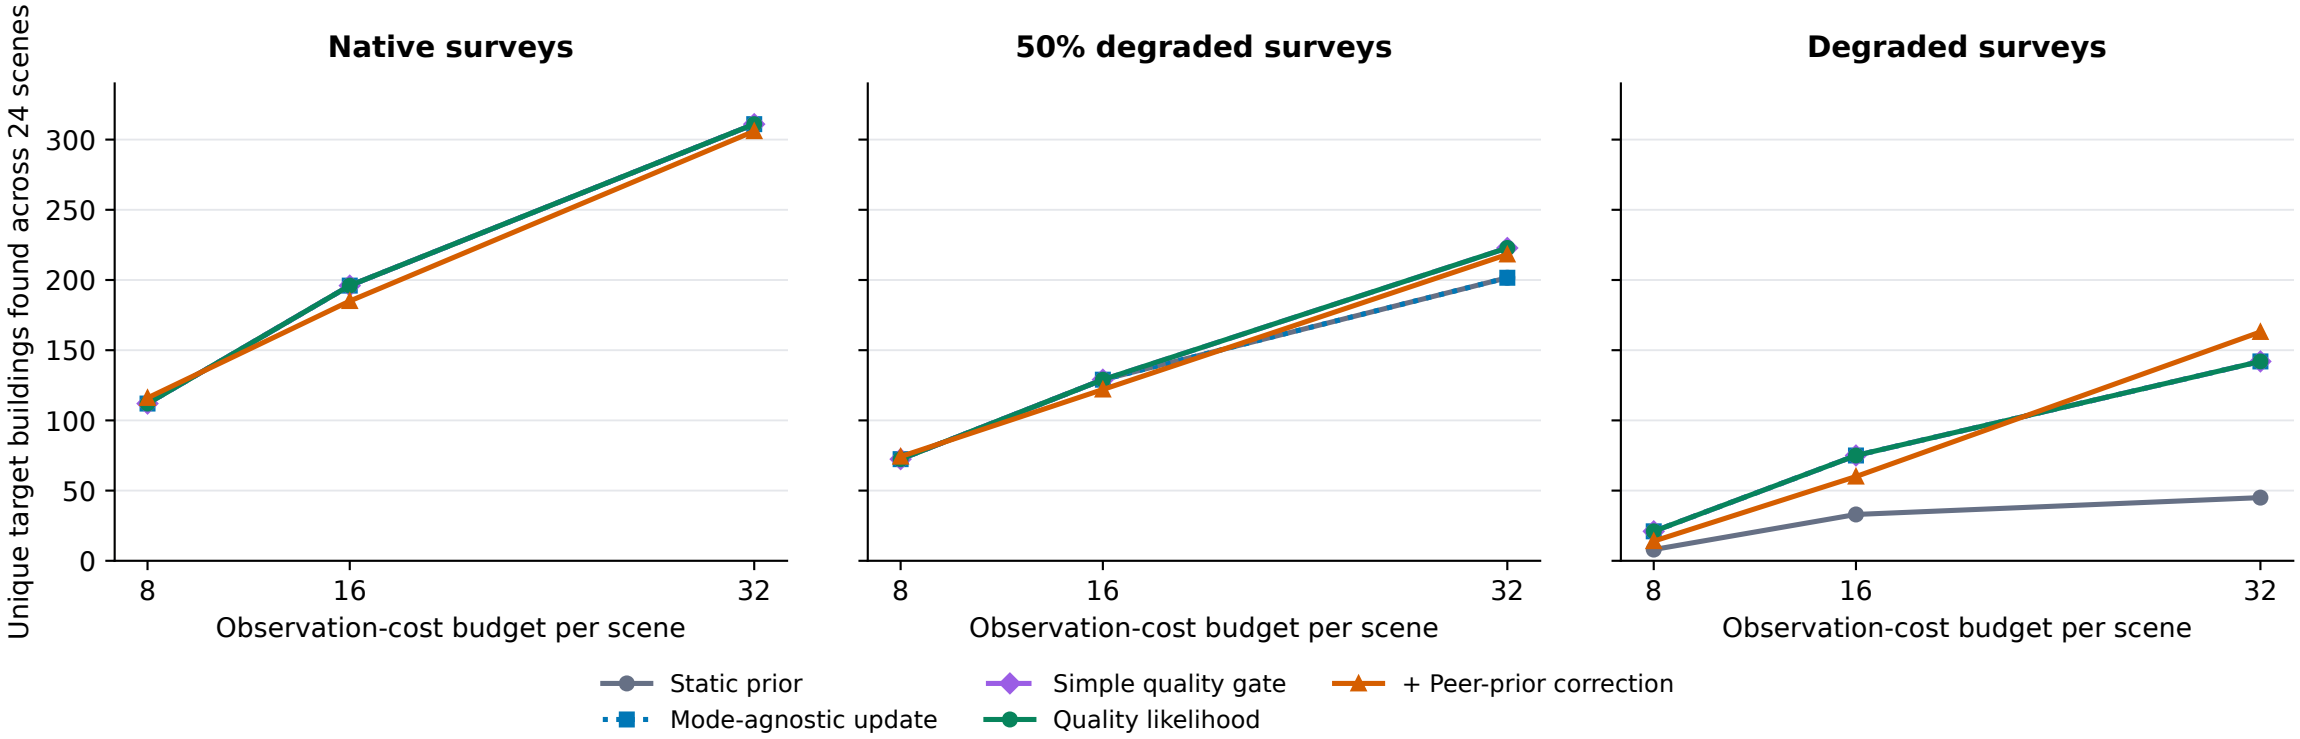
Peer-prior correction changes actions, without consistent discovery gains
Native surveys
50% degraded surveys
Degraded surveys
Unique target buildings found across 24 scenes
0
50
100
150
200
250
300
8
16
32
Observation-cost budget per scene
Static prior
Mode-agnostic update
Simple quality gate
Quality likelihood
+ Peer-prior correction
Calibration only: 24 scenes / 2 events. Mixed results average 5 paired quality maps. Survey cost 1; upgrade cost 2.
Overlapping curves reflect equal outcomes; quality gate and quality likelihood have identical action traces in this run.